

# CSC311 A2

Benson Li, 1007815376

February 15, 2023

## 1 Expected Loss and Bayes Optimality.

### 1.1 (a)

For the policy that keeps every email:

$$\begin{aligned} E[J(y, t = Spam)] &= P(t = Spam) \cdot J(Keep, Spam) + P(t = NonSpam) \cdot J(Keep, NonSpam) \\ &= 0.2 \cdot 1 + 0.8 \cdot 0 \\ &= 0.2 \end{aligned}$$

For the policy that removes every email:

$$\begin{aligned} E[J(y, t = NonSpam)] &= P(t = Spam) \cdot J(Remove, Spam) + P(t = NonSpam) \cdot J(Remove, NonSpam) \\ &= 0.2 \cdot 0 + 0.8 \cdot 100 \\ &= 80 \end{aligned}$$

### 1.2 (b)

We will select the decision  $y$  that minimizes the expected loss:

We denote  $P(t = Spam|x)$  as  $p$

$$\begin{aligned} E[J(y, t)|x] &= P(t = NonSpam|x) \cdot J(y, t = NonSpam) + P(t = Spam|x) \cdot J(y, t = Spam) \\ &= (1 - p) \cdot J(y, t = NonSpam) + p \cdot J(y, t = Spam) \end{aligned}$$

Therefore:

$$\begin{aligned} E[J(y = Keep, t)|x] &= (1 - p) \cdot J(y = Keep, t = NonSpam) + p \cdot J(y = Keep, t = Spam) \\ &= (1 - p) \cdot 0 + p \cdot 1 \\ &= p \end{aligned}$$

$$\begin{aligned} E[J(y = Remove, t)|x] &= (1 - p) \cdot J(y = Remove, t = NonSpam) + p \cdot J(y = Remove, t = Spam) \\ &= (1 - p) \cdot 100 + p \cdot 0 \\ &= 100 - 100p \end{aligned}$$

When  $p \geq \frac{100}{101}$ , we have  $E[J(y = Remove, t)|x] \leq E[J(y = Keep, t)|x]$ , we choose  $y^* = Remove$ .

When  $p < \frac{100}{101}$ , we have  $E[J(y = Keep, t)|x] < E[J(y = Remove, t)|x]$ , we choose  $y^* = Keep$ .

### 1.3 (c)

According to the Baye's Rule:

$$P(t|x_1, x_2) = \frac{P(x_1, x_2|t)P(t)}{P(x_1, x_2)}$$

Then, we are able to calculate that:

$$\begin{aligned} P(t = Spam|x_1 = 0, x_2 = 0) &= \frac{0.45 \cdot 0.2}{0.2 \cdot 0.45 + 0.8 \cdot 0.996} = 0.1015 \\ P(t = Spam|x_1 = 1, x_2 = 0) &= \frac{0.18 \cdot 0.2}{0.2 \cdot 0.18 + 0.8 \cdot 0.002} = 0.9574 \\ P(t = Spam|x_1 = 1, x_2 = 1) &= \frac{0.12 \cdot 0.2}{0.2 \cdot 0.12 + 0.8 \cdot 0} = 1 \\ P(t = Spam|x_1 = 0, x_2 = 1) &= \frac{0.25 \cdot 0.2}{0.2 \cdot 0.25 + 0.8 \cdot 0.002} = 0.9690 \end{aligned}$$

We will select the decision  $y$  that minimizes the expected loss:

$$E[J(y, t)|x_1, x_2] = P(t = NonSpam|x_1, x_2) \cdot J(y, t = NonSpam) + P(t = Spam|x_1, x_2) \cdot J(y, t = Spam)$$

Therefore, according to the given table:

$$\begin{aligned} E[J(y = Keep, t)|x_1 = 0, x_2 = 0] &= 0.8985 \cdot 0 + 0.1015 \cdot 1 \\ &= 0.1015 \end{aligned}$$

$$\begin{aligned} E[J(y = Remove, t)|x_1 = 0, x_2 = 0] &= 0.8985 \cdot 100 + 0.1015 \cdot 0 \\ &= 89.85 \end{aligned}$$

Given  $x_1 = 0 \wedge x_2 = 0$ , since  $0.1015 < 89.85$ , I would keep the email.

$$\begin{aligned} E[J(y = Keep, t)|x_1 = 1, x_2 = 0] &= 0.0426 \cdot 0 + 0.9574 \cdot 1 \\ &= 0.9574 \end{aligned}$$

$$\begin{aligned} E[J(y = Remove, t)|x_1 = 1, x_2 = 0] &= 0.0426 \cdot 100 + 0.9574 \cdot 0 \\ &= 4.26 \end{aligned}$$

Given  $x_1 = 1 \wedge x_2 = 0$ , since  $0.9574 < 4.26$ , I would keep the email.

$$\begin{aligned} E[J(y = Keep, t)|x_1 = 1, x_2 = 1] &= 0 \cdot 0 + 1 \cdot 1 \\ &= 1 \end{aligned}$$

$$\begin{aligned} E[J(y = Remove, t)|x_1 = 1, x_2 = 1] &= 0 \cdot 100 + 1 \cdot 0 \\ &= 0 \end{aligned}$$

Given  $x_1 = 1 \wedge x_2 = 1$ , since  $0 < 1$ , I would remove the email.

$$\begin{aligned} E[J(y = Keep, t)|x_1 = 0, x_2 = 1] &= 0.031 \cdot 0 + 0.969 \cdot 1 \\ &= 0.969 \end{aligned}$$

$$\begin{aligned} E[J(y = Remove, t)|x_1 = 0, x_2 = 1] &= 0.031 \cdot 100 + 0.969 \cdot 0 \\ &= 3.1 \end{aligned}$$

Given  $x_1 = 0 \wedge x_2 = 1$ , since  $0.969 < 3.1$ , I would keep the email.

#### 1.4 (d)

$$P(x_1 = 0, x_2 = 0) = 0.2 \cdot 0.45 + 0.8 \cdot 0.996 = 0.8868$$

$$P(x_1 = 1, x_2 = 0) = 0.2 \cdot 0.18 + 0.8 \cdot 0.002 = 0.0376$$

$$P(x_1 = 1, x_2 = 1) = 0.2 \cdot 0.12 + 0.8 \cdot 0 = 0.024$$

$$P(x_1 = 0, x_2 = 1) = 0.2 \cdot 0.25 + 0.8 \cdot 0.002 = 0.0516$$

The expectation loss is:

$$\begin{aligned} E[J(y^*, t)] &= \sum_{x_1, x_2} P(x_1, x_2) E[J(y^*, t) | x_1, x_2] \\ &= 0.8868 \cdot 0.1015 + 0.0376 \cdot 0.9574 + 0 \cdot 0.024 + 0.0516 \cdot 0.969 \\ &= 0.176 \end{aligned}$$

## 2 Feature Maps.

### 2.1 (a)

Assume this data-set is linear separable:

Let  $x_0 = 1, w_0 = b(\text{bias}), X$  is design matrix.

Then  $\bar{y} = w^T X$ , we will make decision based on the boundary 0.

Therefore, we got following equations:

$$w_0 + w_1(-1) + w_2(0) \geq 0 \quad (1)$$

$$w_0 + w_1(1) + w_2(2) < 0 \quad (2)$$

$$w_0 + w_1(2) + w_2(3) \geq 0 \quad (3)$$

We will try to find solution for these equations:

From (3) - (2), we get  $w_1 + w_2 > 0$ , we denote this as (4)

We plus  $2 * (4)$  to (1), we then get that  $w_0 + w_1 + 2w_2 > 0$ .

This contradicts with equation(2), which proves that there is no solution for these equations.

Therefore, this data-set is not linearly separable.

### 2.2 (b)

Let  $x_0 = 1, X$  is design matrix.

We will make decision based on the boundary 0.

Therefore, we got following equations:

$$w_0 + w_1(-1) + w_2(0) + w_3(0) \geq 0 \quad (1)$$

$$w_0 + w_1(1) + w_2(2) + w_3(4) < 0 \quad (2)$$

$$w_0 + w_1(2) + w_2(3) + w_3(9) \geq 0 \quad (3)$$

Equivalently:

$$w_0 - w_1 \geq 0 \quad (1)$$

$$w_0 + w_1 + 2w_2 + 4w_3 < 0 \quad (2)$$

$$w_0 + 2w_1 + 3w_2 + 9w_3 \geq 0 \quad (3)$$

These are the constraints that  $w_0, w_1, w_2, w_3$  need to satisfy.

Even though I don't need to solve  $w_0, w_1, w_2, w_3$ , I list my solution below just for practise.

From (1), we can have  $w_0 = 2, w_1 = 1$ .

We plug this in (2), we need to satisfy  $3 + 2w_2 + 4w_3 < 0$ , we can have  $w_2 = -6, w_3 = 2$ .

We plug this in (3), we check that  $2 + 2 + (-18) + 18 = 4 \geq 0$ .

Therefore, we have the following solution:

$$w_0 = 2$$

$$w_1 = 1$$

$$w_2 = -6$$

$$w_3 = 2$$

### 3 kNN vs. Logistic Regression

#### 3.1 (a)

##### 3.1.1 (i)

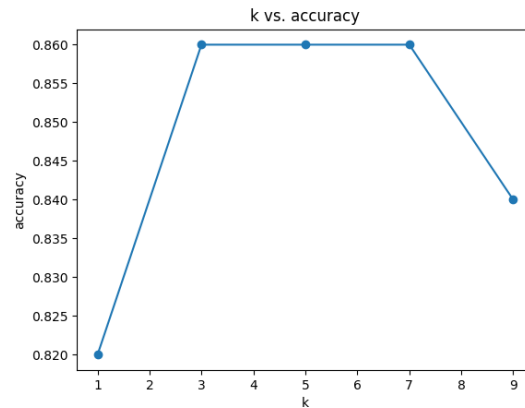


Figure 1: The plot for accuracy vs. k

##### 3.1.2 (ii)

```
Python Console x run_knn x
/usr/local/bin/python3.10 /Applications/PyCharm.app/Contents/plugins/python/helpers/pydev/pydevconsole.py
import sys; print('Python %s on %s' % (sys.version, sys.platform))
sys.path.extend(['/Users/qiu/Desktop/Mac/Github/CSC311'])

Python 3.10.9 (v3.10.9:1dd9be6584, Dec 6 2022, 14:37:36) [Clang 13.0.0 (clang-1300.0.29.30)]
+ I choose k = 5, since it's the median value for highest accuracies
Accuracy on validation set: 0.86
Accuracy on test set: 0.94
=====
Try k = 7 (5 + 2)
Accuracy on validation set: 0.86
Accuracy on test set: 0.94
=====
Try k = 3 (5 - 2)
Accuracy on validation set: 0.86
Accuracy on test set: 0.94
=====
As we change the value of k, we can see that the accuracy don' have a big change in this question.
However, we expect the accuracy to decrease when we change the value of k.
When k is smaller than 5, the model would more likely to underfit the data.
When k is larger than 5, the model would more likely to overfit the data.

In [3]: |
```

Figure 2: Test and validation accuracies change as we change the value of k

As we change the value of k, we can see that the accuracy don' have a big change in this question. However, we expect the accuracy to decrease when we change the value of k. When k is smaller than 5, the model would more likely to underfit the data. When k is larger than 5, the model would more likely to overfit the data.

## 3.2 (b)

### 3.2.1 (i)

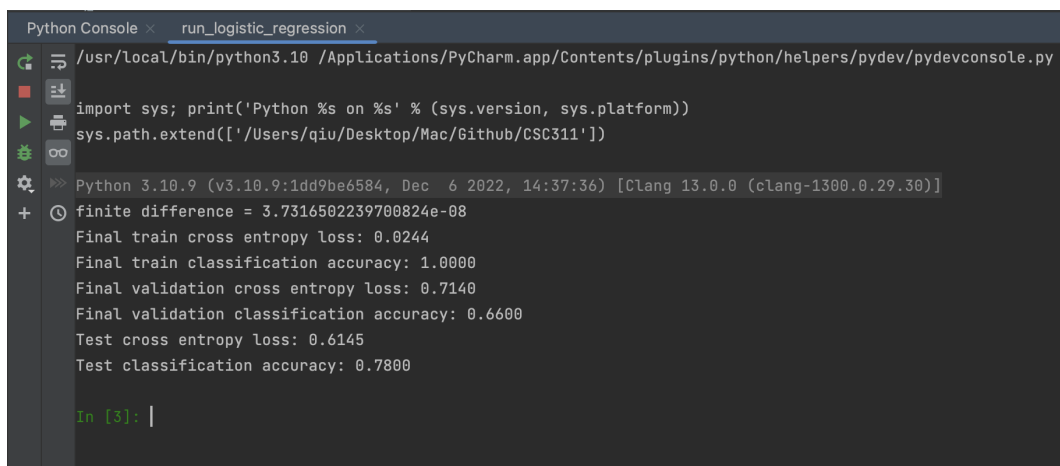
Please refer to python files.

### 3.2.2 (ii)

The following result (for question (ii) and (iii)) is collected under hyperparameters (best hyperparameter settings):

**learning\_rate:** 0.01;  
**weight\_regularization** = 0;  
**num\_iterations:** 1000

The output result for small training set:



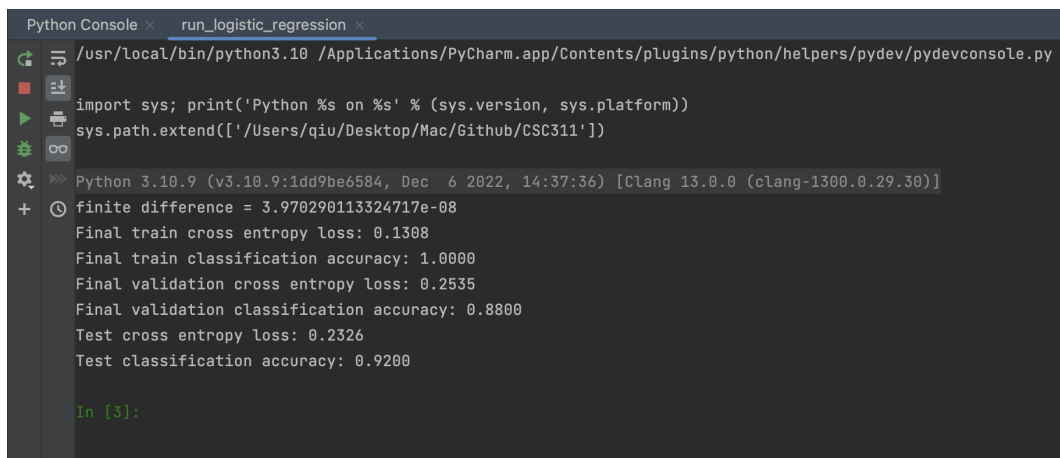
```
Python Console x run_logistic_regression x
/usr/local/bin/python3.10 /Applications/PyCharm.app/Contents/plugins/python/helpers/pydev/pydevconsole.py
import sys; print('Python %s on %s' % (sys.version, sys.platform))
sys.path.extend(['/Users/qiu/Desktop/Mac/Github/CSC311'])

>> Python 3.10.9 (v3.10.9:1dd9be6584, Dec 6 2022, 14:37:36) [Clang 13.0.0 (clang-1300.0.29.30)]
+ finite difference = 3.7316502239700824e-08
  Final train cross entropy loss: 0.0244
  Final train classification accuracy: 1.0000
  Final validation cross entropy loss: 0.7140
  Final validation classification accuracy: 0.6600
  Test cross entropy loss: 0.6145
  Test classification accuracy: 0.7800

In [3]: |
```

Figure 3: Gradient checking and Test result for small dataset

The output result for large training set:



```
Python Console x run_logistic_regression x
/usr/local/bin/python3.10 /Applications/PyCharm.app/Contents/plugins/python/helpers/pydev/pydevconsole.py
import sys; print('Python %s on %s' % (sys.version, sys.platform))
sys.path.extend(['/Users/qiu/Desktop/Mac/Github/CSC311'])

>> Python 3.10.9 (v3.10.9:1dd9be6584, Dec 6 2022, 14:37:36) [Clang 13.0.0 (clang-1300.0.29.30)]
+ finite difference = 3.970290113324717e-08
  Final train cross entropy loss: 0.1308
  Final train classification accuracy: 1.0000
  Final validation cross entropy loss: 0.2535
  Final validation classification accuracy: 0.8800
  Test cross entropy loss: 0.2326
  Test classification accuracy: 0.9200

In [3]:
```

Figure 4: Gradient checking and Test result for large dataset

### 3.2.3 (iii)

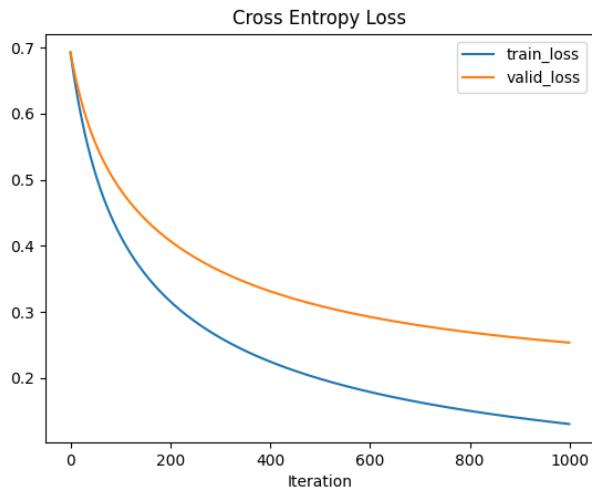


Figure 5: Cross Entropy Loss

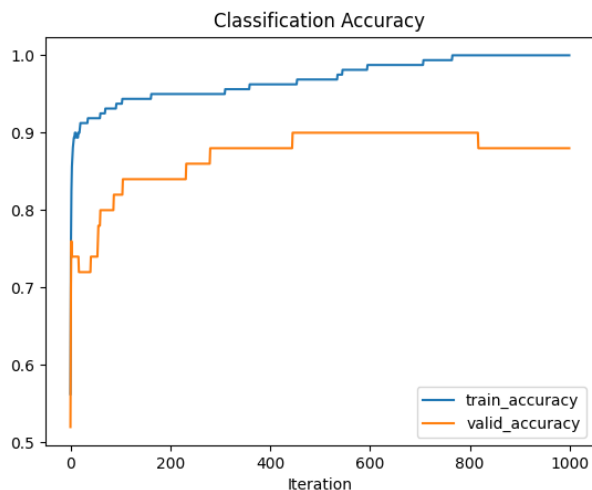


Figure 6: Classification Accuracy

### 3.2.4 (iv)

We run our program for three times with the same hyper-parameters on the large dataset and got the following result.

Result				
	Train_Loss	Train_Accuracy	Val_Loss	Val_Accuracy
1st	0.0244	1.0000	0.7140	0.6600
2nd	0.1307	1.0000	0.2538	0.8800
3rd	0.1308	1.0000	0.2536	0.8800

To choose the best hyper-parameters, we can run the code multiple times and calculate the average validation accuracy error. This will help to ensure that the results are not simply due to chance.