

Design Theory for Relational Databases

Author: Diane Horton

Originally based on slides by Jeff Ullman

Introduction

- There are always many different schemas for a given set of data.
- E.g., you could combine or divide tables.
- How do you pick a schema?
Which is better?
What does “better” mean?
- Fortunately, there are some principles to guide us.

Database Design Theory

- It allows us to improve a schema systematically.
- General idea:
 - Express constraints on the relationships between attributes
 - Use these to decompose the relations
- Ultimately, get a schema that is in a “normal form” that guarantees good properties.
- “Normal” in the sense of conforming to a standard.
- The process of converting a schema to a normal form is called **normalization**.

Part I: Functional Dependency Theory

A poorly designed table

part	manufacturer	manAddress	seller	sellerAddress	price
1983	Hammers `R Us	99 Pinecrest	ABC	1229 Bloor W	5.59
8624	Lee Valley	102 Vaughn	ABC	1229 Bloor W	23.99
9141	Hammers `R Us	99 Pinecrest	ABC	1229 Bloor W	12.50
1983	Hammers `R Us	99 Pinecrest	Walmart	5289 St Clair W	4.99

Perhaps:

- Every part has 1 manufacturer
- Every manufacturer has 1 address
- Every seller has 1 address

If so, instances will have redundant data.

Principle: Avoid redundancy

- Redundant data can lead to anomalies.

part	manufacturer	manAddress	seller	sellerAddress	price
1983	Hammers `R Us	99 Pinecrest	ABC	1229 Bloor W	5.59
8624	Lee Valley	102 Vaughn	ABC	1229 Bloor W	23.99
9141	Hammers `R Us	99 Pinecrest	ABC	1229 Bloor W	12.50
1983	Hammers `R Us	99 Pinecrest	Walmart	5289 St Clair W	4.99

- **Update anomaly:** if Hammers `R Us moves and we update only one tuple, the data is inconsistent!

part	manufacturer	manAddress	seller	sellerAddress	price
1983	Hammers `R Us	99 Pinecrest	ABC	1229 Bloor W	5.59
8624	Lee Valley	102 Vaughn	ABC	1229 Bloor W	23.99
9141	Hammers `R Us	99 Pinecrest	ABC	1229 Bloor W	12.50
1983	Hammers `R Us	99 Pinecrest	Walmart	5289 St Clair W	4.99

- **Deletion anomaly:** if ABC stops selling part 8624 and Lee Valley makes only that one part, we lose track of its address!

Definition of FD

- Suppose R is a relation, and X and Y are subsets of the attributes of R .
- $X \rightarrow Y$ asserts that:
 - If two tuples agree on all the attributes in set X , they must also agree on all the attributes in set Y .
- We say that " $X \rightarrow Y$ holds in R ", or " X functionally determines Y ."
- An FD constrains what can go in a relation.

More formally...

$A \rightarrow B$ means:

\forall tuples $t_1, t_2,$

$$(t_1[A] = t_2[A]) \Rightarrow (t_1[B] = t_2[B])$$

Or equivalently:

$\neg \exists$ tuples t_1, t_2 such that

$$(t_1[A] = t_2[A]) \wedge (t_1[B] \neq t_2[B])$$

Generalization to multiple attributes

$A_1 A_2 \dots A_m \rightarrow B_1 B_2 \dots B_n$ means:

\forall tuples t_1, t_2 ,

$(t_1[A_1] = t_2[A_1] \wedge \dots \wedge t_1[A_m] = t_2[A_m]) \Rightarrow$

$(t_1[B_1] = t_2[B_1] \wedge \dots \wedge t_1[B_n] = t_2[B_n])$

Or equivalently:

$\neg \exists$ tuples t_1, t_2 such that

$(t_1[A_1] = t_2[A_1] \wedge \dots \wedge t_1[A_m] = t_2[A_m]) \wedge$

$\neg (t_1[B_1] = t_2[B_1] \wedge \dots \wedge t_1[B_n] = t_2[B_n])$

FDs in our parts table

part	manufacturer	manAddress	seller	sellerAddress	price
1983	Hammers `R Us	99 Pinecrest	ABC	1229 Bloor W	5.59

etc.

We can now express relationships as FDs

- Every part has 1 manufacturer becomes
 $\text{part} \rightarrow \text{manufacturer}$
- Every manufacture has 1 address becomes
 $\text{manufacturer} \rightarrow \text{manAddress}$
- Every seller has 1 address becomes
 $\text{seller} \rightarrow \text{sellerAddress}$

Why “functional dependency”?

- $X \rightarrow Y$ is a “dependency” because
the value of Y depends on the value of X .
- It is “functional” because
there is a mathematical function that takes a
value for X and gives a *unique* value for Y .
(It’s not a typical function; just a lookup.)

Equivalent sets of FDs

- When we write a set of FDs, we mean that all of them hold.
- We can very often rewrite sets of FDs in equivalent ways.
- When we say S_1 is equivalent to S_2 we mean that:
 S_1 holds in a relation iff S_2 does.

Splitting rules for FDs

- Can we split the RHS of an FD and get multiple, equivalent FDs?
- Can we split the LHS of an FD and get multiple, equivalent FDs?

Coincidence or FD?

- An FD is an assertion about *every* instance of the relation.
- You can't know it holds just by looking at one instance.
- You must use knowledge of the domain to determine whether an FD holds.

FDs are closely related to keys

- Suppose K is a set of attributes for relation R .
- Our old definition of superkey:
a set of attributes for which no two rows can have the same values.
- A claim about FDs:
 K is a **superkey** for R iff
 K functionally determines all of R .

FDs are a generalization of keys

- Superkey:
 $X \rightarrow R$
Every attribute
- Functional dependency:
 $X \rightarrow Y$
Not necessarily every attribute
- An FD can be more subtle.

Inferring FDs

- Given a set of FDs, we can often infer further FDs.
- This will be handy when we apply FDs to the problem of database design.
- Big task: given a set of FDs,
infer *every* other FD that must also hold.
- Simpler task: given a set of FDs,
check whether *a given* FD must also hold.

Examples

- If $A \rightarrow B$ and $B \rightarrow C$ hold,
must $A \rightarrow C$ hold?
- If $A \rightarrow H$, $C \rightarrow F$, and $FG \rightarrow AD$ hold,
must $FA \rightarrow D$ hold?
must $CG \rightarrow FH$ hold?
- If $H \rightarrow GD$, $HD \rightarrow CE$, and $BD \rightarrow A$ hold,
must $EH \rightarrow C$ hold?
- Aside: we are not generating new FDs,
but testing a specific possible one.

Method 1: Prove an FD follows using first principles

- You can prove it by referring back to
 - The FDs that you know hold, and
 - The definition of functional dependency.

Method 2: Prove an FD follows using the Closure Test

- Assume you know the values of the LHS attributes, and figure out everything else that is determined.
- If it includes the RHS attributes, then you know that $LHS \rightarrow RHS$
- This is called the closure test.

*Y is a set of attributes, S is a set of FDs.
Return the closure of Y under S.*

Attribute_closure(Y, S):

Initialize Y^+ to Y


Repeat until no more changes occur:

 If there is an FD LHS \rightarrow RHS in S
 such that LHS is in Y^+ :

 Add RHS to Y^+

Return Y^+

Visualizing attribute closure



If LHS is in Y^+ and $\text{LHS} \rightarrow \text{RHS}$ holds, we can add RHS to Y^+

*S is a set of FDs; LHS \rightarrow RHS is a single FD.
Return true iff LHS \rightarrow RHS follows from S.*

FD_follows(S, LHS \rightarrow RHS):

$Y^+ = \text{Attribute_closure}(\text{LHS}, S)$

return (RHS is in Y^+)

Projecting FDs

- Later, we will learn how to **normalize** a schema by **decomposing** relations.

This is the whole point of this theory.

- We will need to know what FDs hold in the new, smaller, relations.

We must **project** our FDs onto the attributes of our new relations.

S is a set of FDs; L is a set of attributes.

Return the projection of S onto L:

all FDs that follow from S and involve only attributes from L.

Project(S, L):

Initialize T to {}.

For each subset X of L:

Compute X^+ *Close X and see what we get.*

For every attribute A in X^+ :

If A is in L: *$X \rightarrow A$ is only relevant if A is in L (we know X is).*

add $X \rightarrow A$ to T.

Return T.

A few speed-ups

- No need to add $X \rightarrow A$ if A is in X itself. It's a trivial FD.
- These subsets of X won't yield anything, so no need to compute their closures:
 - the empty set
 - the set of all attributes
- Neither are big savings, but ...

A big speed-up

- If we find that $X^+ = \text{all attributes}$, we can ignore any superset of X .

It can only give use “weaker” FDs
(i.e., FDs with more on the LHS).

- This is a big time saver!

Projection is expensive

- Even with these speed-ups, projection is still expensive.
- Suppose R_1 has n attributes.
How many subsets of R_1 are there?

Minimal Basis

- We saw earlier that we can very often rewrite sets of FDs in equivalent ways.
- Example: $S_1 = \{A \rightarrow BC\}$ is equivalent to $S_2 = \{A \rightarrow B, A \rightarrow C\}$.
- Given a set of FDs S , we may want to find a **minimal basis**: A set of FDs that is equivalent, but has
 - no redundant FDs, and
 - no FDs with unnecessary attributes on the LHS.

S is a set of FDs. Return a minimal basis for S.

Minimal_basis(S):

1. Split the RHS of each FD
2. For each FD $X \rightarrow Y$ where $|X| \geq 2$:
If you can remove an attribute from X
and get an FD that follows from S:
Do so! (It's a stronger FD.)
3. For each FD f :
If $S - \{f\}$ implies f :
Remove f from S.

Some comments on minimal basis

- Often there are multiple possible results.
Depends on the order in which you consider the possible simplifications.
- After you identify a redundant FD, you must not use it when computing subsequent closures.

... and perhaps less intuitive

- When you are computing closures to decide whether the LHS of an FD
 $X \rightarrow Y$
can be simplified, continue to use that FD.
- You must do steps (2) and (3) in that order.
Otherwise, must repeat until no changes occur.