

FROM E/R MODEL TO DATABASE SCHEMA

Steps

- *Restructure* the ER schema to improve it, based on criteria
- *Translate* the schema into the relational model
- *Add missing constraints* to the schema

1. RESTRUCTURING AN E/R MODEL

Restructuring Overview

Input: E/R Schema

Output: Restructured E/R Schema


It includes (not necessarily in this order):

- a. Analysis of redundancies
- b. Choosing entity set vs attribute
- c. Limiting the use of weak entity sets
- d. Settling on keys
- e. Creating entity sets to replace attributes with cardinality greater than one

1a. Analysis of redundancies


Example: no redundancy

It is not redundant to have Name twice.




Example: redundancy

What is redundant here?



Example: redundancy

What is redundant here?



1b. Entity sets vs attributes



Overall, Prefer Attributes

- An entity set should satisfy at least one of the following conditions:
 - It is more than the name of something; it has at least one non-key attribute, or
 - It is the “many” in a many-one or many-many relationship.
- Rules of thumb
 - A “thing” in its own right => Entity Set
 - A “detail” about some other “thing” => Attribute

This is just about avoiding redundancy



E.S. vs. attributes: examples

Domain fact change: **A part can have more than one manufacturer ...**





E.S. vs. attributes: examples

Domain fact change: **Not representing Manufacturer address ...**





E.S. vs. attributes: examples

New domain



E.S. vs. attributes: examples

Domain fact change: **A mentor can have more than one mentee ...**



1c. Limiting weak entity sets

When to use weak entity sets?

- The usual reason is that there is no global authority capable of creating unique ID's
- **Example:** it is unlikely that there could be an agreement to assign unique student numbers across all students in the world
- It may seem the only choice is a weak entity set.
- It is usually better to create unique IDs
 - Social insurance number, automobile VIN, etc.

1d. Settling on keys

Settling on keys

- Make sure that every entity set has a key
- Keep in mind:
 - Attributes with null values cannot be part of primary keys
 - Internal keys are preferable to external ones
 - A key that is used by many operations to access instances of an entity is preferable to others
 - A key with one/few attributes is preferable
 - An integer key is preferable

Avoid multi-attribute and string keys


- They are wasteful
 - e.g. Movies(title, year, ...): 2 attributes in key, requires ~16 bytes
 - Number of movies ever made ($\ll 2^{32}$) can be distinguished with 4 bytes
 - => Having an integer movieID key saves 75% space and a lot of typing
- They break encapsulation
 - e.g. Patient(firstName, lastName, phone, ...)
 - Security/privacy hole
 - => Integer patientID prevents information leaks
- They are brittle (nasty interaction of above two points)
 - Name or phone number change? Parent and child with same name?
 - Patient with no phone? Two movies with same title and year?
 - => Internal ID always exist, are immutable, unique

Also: computers are really good at integers!

1e. Creating entity sets to
replace attributes with
cardinality greater than one

Attributes with cardinality > 1

- The relational model doesn't allow multi-valued attributes. We must convert these to entity sets.



2. TRANSLATING AN E/R MODEL INTO A DB SCHEMA

Translation into a Logical Schema

Input: E/R Schema

Output: Relational Schema

- Starting from an E/R schema, an equivalent relational schema is constructed
 - “equivalent”: a schema capable of representing the same information
- A good translation should also:
 - not allow redundancy
 - not invite unnecessary null values

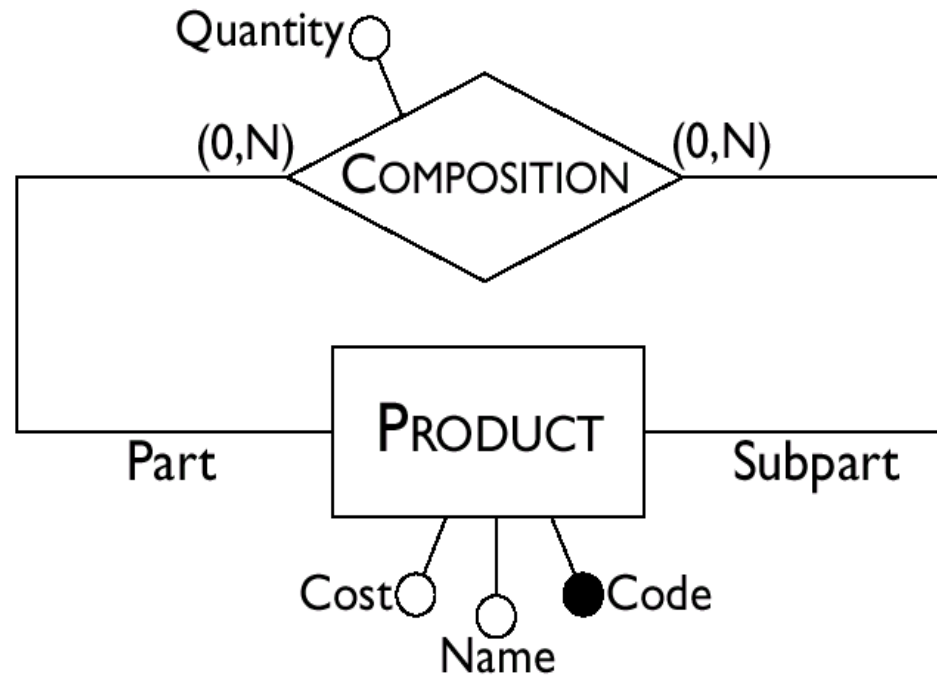
The general idea

- Each entity set becomes a relation.
Its attributes are
 - the attributes of the entity set.
- Each relationship becomes a relation.
Its attributes are
 - the keys of the entity sets that it connects, plus
 - the attributes of the relationship itself.
- We'll see opportunities to simplify.

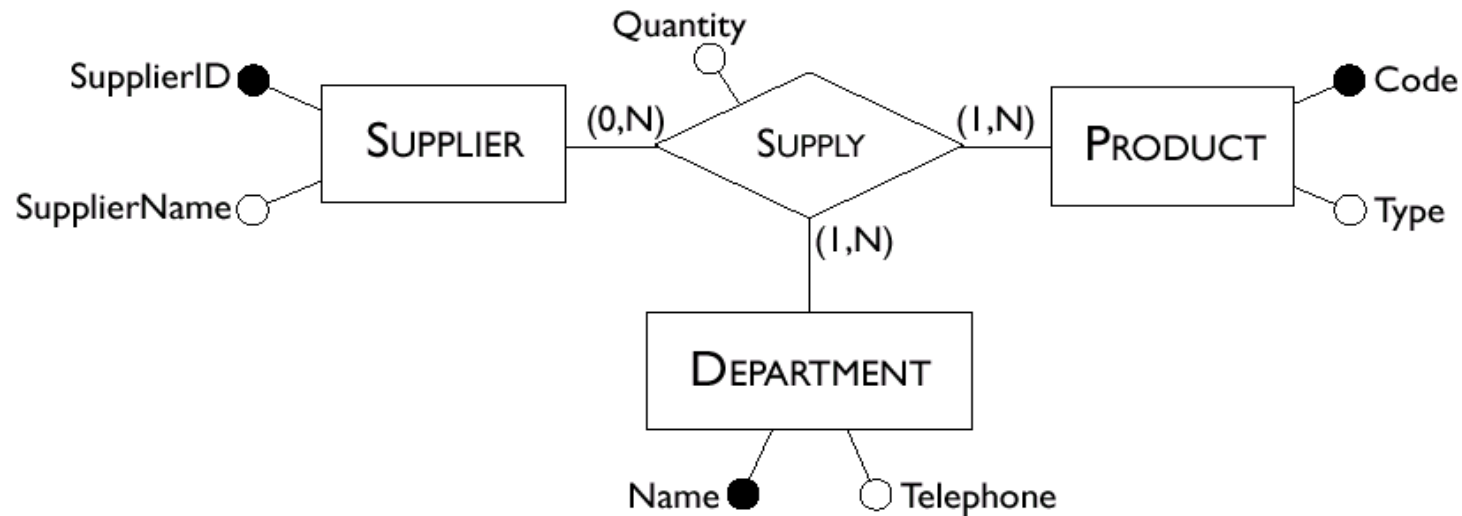
Many-to-Many Binary Relationships



Many-to-Many Recursive Relationships



Many-to-Many Ternary Relationships




Simplifications

- These straight translations are acceptable.
- But we can often simplify.

One-to-Many Relationships

with mandatory participation on the “one” side



One-to-One Relationships

with mandatory participation for both



One-to-One Relationships

with mandatory participation for both



The standard translation has 3 relations (what is key of management)?

One-to-One Relationships

with optional participation for one



Summary of Types of Relationship

- many-to-many (binary or ternary)
- one-to-many
 - mandatory: (1,1) on the “one” side
 - optional: (0,1) on the “one” side
- one-to-one
 - both mandatory: (1,1) on both sides
 - one mandatory, one optional:
(1,1) on one side and (0,1) on other side
 - both optional: (0,1) on both sides

3. FINAL CONSIDERATIONS

During or after the translation, we should be sure to:

- Express the foreign-key constraints
- Consider better names for the referring attributes
- Express the "max 1" constraints
- Express the "min 1" constraints

Names for Referring Attributes

Example:



Employee(Number, Surname, Salary)

Project(Code, Name, Budget)

Participation(Number, Code, StartDate)

Participation[Number] \subseteq Employee[Number]

Participation[Code] \subseteq Project[Code]

What about “max 1” constraints?



Our simplified v1 (with Management info moved over to Head):

Head(Number, Name, Salary, Department, StartDate)

Department(Name, Telephone, Branch)

$\text{Head}[\text{Department}] \subseteq \text{Department}[\text{Name}]$

Did we enforce that every head has at most one department?

What about “max 1” constraints?



Our simplified v1 (with Management info moved over to Head):


Head(Number, Name, Salary, Department, StartDate)

Department(Name, Telephone, Branch)

$\text{Head}[\text{Department}] \subseteq \text{Department}[\text{Name}]$

Did we enforce that every department has at most one head?

What about “min 1” constraints?



Our simplified translation:


Player(Surname, DOB, Position, TeamName, Salary)

Team(Name, Town, TeamColours)

Player[TeamName] \subseteq Team[Name]

Did we enforce that every player must have a team?

What about “min 1” constraints?



Our standard (non simplified) translation:

Player(Surname, DOB, Position)

Team(Name, Town, TeamColours)

Contract(PlayerSurname, PlayerDOB, Team, Salary)

Contract[PlayerSurname, PlayerDOB] \subseteq Player[Surname, DOB]

Contract[Team] \subseteq Team[Name]

Did we enforce that every player must have a team?

What about “min 1” constraints?



Our simplified translation:

Player(Surname, DOB, Position, TeamName, Salary)

Team(Name, Town, TeamColours)

Player[TeamName] \subseteq Team[Name]

Does this enforce that every team must have a player?

What about “min 1” constraints?

In our translation from E/R to a relational schema, we expressed the constraints using relational notation.

But in SQL, only some of them can be written as foreign keys.

- They must refer to attributes that are primary key or unique.

Will the schema be “good”?

- If we use this process, will the schema we get be a good one?
- The process should ensure that (a) the data can be represented, (b) there is no redundancy, and (c) most constraints are enforced.
- But only with respect to what’s in the E/R diagram.
- Crucial thing we are missing: functional dependencies. (We only have keys, not other FDs.)
- So a good design process includes normalization at the end.

Redundancy can be *desirable*

- **Disadvantages** of redundancy:
 - More storage (but usually at negligible cost)
 - Additional operations to keep the data consistent
- **Advantages** of redundancy:
 - Speed: Fewer accesses necessary to obtain information
- So a good design process includes considering whether to allow some redundancy.
- How to decide to maintain or eliminate a redundancy?
Examine:
 - the speedup in operations made possible by the redundant information
 - the relative frequency of those operations
 - the storage needed for the redundant informations